

Regression Analysis of Regularization and Dimension Reduction Techniques

Ambachow Kahsay

2025-10-16

Contents

Introduction

This report explores the California Housing dataset, performs Exploratory Data Analysis (EDA), and builds a Multiple Linear Regression model to predict median house values.

```
library(recipes)
```

0.0.0.1 Libraries

```
## Loading required package: dplyr
## Warning: package 'dplyr' was built under R version 4.3.3
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
##
## Attaching package: 'recipes'
## The following object is masked from 'package:stats':
##
##   step
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.3.3
## Loading required package: ggplot2
## Loading required package: lattice
```

```
library(magrittr)
```

```
library(magrittr)
```

```
library(caret)
library(magrittr)
library(glmnet)
```

```
## Loading required package: Matrix
## Loaded glmnet 4.1-10
```

```
# Load the dataset using the absolute path
housing <- read.csv("C:/Users/Admin/OneDrive - United States International University (USIU)/Documents/

# View the first few rows of the data
head(housing)
```

0.0.0.1 Load the Dataset

```
## longitude latitude housing_median_age total_rooms total_bedrooms population
## 1 -122.23 37.88 41 880 129 322
## 2 -122.22 37.86 21 7099 1106 2401
## 3 -122.24 37.85 52 1467 190 496
## 4 -122.25 37.85 52 1274 235 558
## 5 -122.25 37.85 52 1627 280 565
## 6 -122.25 37.85 52 919 213 413
## households median_income median_house_value ocean_proximity
## 1 126 8.3252 452600 NEAR BAY
## 2 1138 8.3014 358500 NEAR BAY
## 3 177 7.2574 352100 NEAR BAY
## 4 219 5.6431 341300 NEAR BAY
## 5 259 3.8462 342200 NEAR BAY
## 6 193 4.0368 269700 NEAR BAY
```

```
# Check the structure of the dataset
str(housing)
```

0.0.0.2 Inspect Data

```
## 'data.frame': 20640 obs. of 10 variables:
## $ longitude : num -122 -122 -122 -122 -122 ...
## $ latitude : num 37.9 37.9 37.9 37.9 37.9 ...
## $ housing_median_age: num 41 21 52 52 52 52 52 52 42 52 ...
## $ total_rooms : num 880 7099 1467 1274 1627 ...
## $ total_bedrooms : num 129 1106 190 235 280 ...
## $ population : num 322 2401 496 558 565 ...
## $ households : num 126 1138 177 219 259 ...
## $ median_income : num 8.33 8.3 7.26 5.64 3.85 ...
## $ median_house_value: num 452600 358500 352100 341300 342200 ...
## $ ocean_proximity : chr "NEAR BAY" "NEAR BAY" "NEAR BAY" "NEAR BAY" ...
```

```
housing <- na.omit(housing)
head(housing)
```

0.0.0.3 Handle Missing Values

```
##   longitude latitude housing_median_age total_rooms total_bedrooms population
## 1   -122.23   37.88             41           880           129           322
## 2   -122.22   37.86             21          7099          1106          2401
## 3   -122.24   37.85             52          1467           190           496
## 4   -122.25   37.85             52          1274           235           558
## 5   -122.25   37.85             52          1627           280           565
## 6   -122.25   37.85             52           919           213           413
##   households median_income median_house_value ocean_proximity
## 1         126         8.3252         452600      NEAR BAY
## 2        1138         8.3014         358500      NEAR BAY
## 3         177         7.2574         352100      NEAR BAY
## 4         219         5.6431         341300      NEAR BAY
## 5         259         3.8462         342200      NEAR BAY
## 6         193         4.0368         269700      NEAR BAY
```

```
housing <- dummyVars(" ~ .", data = housing) %>%
  predict(newdata = housing) %>%
  as.data.frame()
```

```
dummies <- dummyVars(" ~ .", data = housing)
housing <- predict(dummies, newdata = housing)
housing <- as.data.frame(housing)
```

```
str(housing)
```

0.0.0.4 Encode Categorical Variable

```
## 'data.frame':   20433 obs. of  14 variables:
## $ longitude      : num  -122 -122 -122 -122 -122 ...
## $ latitude       : num  37.9 37.9 37.9 37.9 37.9 ...
## $ housing_median_age : num  41 21 52 52 52 52 52 52 42 52 ...
## $ total_rooms    : num  880 7099 1467 1274 1627 ...
## $ total_bedrooms : num  129 1106 190 235 280 ...
## $ population     : num  322 2401 496 558 565 ...
## $ households     : num  126 1138 177 219 259 ...
## $ median_income  : num  8.33 8.3 7.26 5.64 3.85 ...
## $ median_house_value : num  452600 358500 352100 341300 342200 ...
## $ `ocean_proximity<1H OCEAN` : num  0 0 0 0 0 0 0 0 0 0 ...
## $ ocean_proximityINLAND : num  0 0 0 0 0 0 0 0 0 0 ...
## $ ocean_proximityISLAND : num  0 0 0 0 0 0 0 0 0 0 ...
## $ `ocean_proximityNEAR BAY` : num  1 1 1 1 1 1 1 1 1 1 ...
## $ `ocean_proximityNEAR OCEAN` : num  0 0 0 0 0 0 0 0 0 0 ...
```

```
summary(housing)
```

```
##   longitude      latitude housing_median_age total_rooms
## Min.   : -124.3   Min.   :32.54   Min.    : 1.00   Min.    :    2
## 1st Qu.: -121.8   1st Qu.:33.93   1st Qu.:18.00   1st Qu.: 1450
## Median : -118.5   Median :34.26   Median :29.00   Median : 2127
## Mean   : -119.6   Mean    :35.63   Mean    :28.63   Mean    : 2636
## 3rd Qu.: -118.0   3rd Qu.:37.72   3rd Qu.:37.00   3rd Qu.: 3143
## Max.   : -114.3   Max.    :41.95   Max.    :52.00   Max.    :39320
## total_bedrooms  population      households median_income
```

```
## Min. : 1.0 Min. : 3 Min. : 1.0 Min. : 0.4999
## 1st Qu.: 296.0 1st Qu.: 787 1st Qu.: 280.0 1st Qu.: 2.5637
## Median : 435.0 Median : 1166 Median : 409.0 Median : 3.5365
## Mean : 537.9 Mean : 1425 Mean : 499.4 Mean : 3.8712
## 3rd Qu.: 647.0 3rd Qu.: 1722 3rd Qu.: 604.0 3rd Qu.: 4.7440
## Max. : 6445.0 Max. : 35682 Max. : 6082.0 Max. : 15.0001
## median_house_value `ocean_proximity<1H OCEAN` ocean_proximityINLAND
## Min. : 14999 Min. : 0.0000 Min. : 0.0000
## 1st Qu.: 119500 1st Qu.: 0.0000 1st Qu.: 0.0000
## Median : 179700 Median : 0.0000 Median : 0.0000
## Mean : 206864 Mean : 0.4421 Mean : 0.3179
## 3rd Qu.: 264700 3rd Qu.: 1.0000 3rd Qu.: 1.0000
## Max. : 500001 Max. : 1.0000 Max. : 1.0000
## ocean_proximityISLAND `ocean_proximityNEAR BAY` `ocean_proximityNEAR OCEAN`
## Min. : 0.0000000 Min. : 0.0000 Min. : 0.0000
## 1st Qu.: 0.0000000 1st Qu.: 0.0000 1st Qu.: 0.0000
## Median : 0.0000000 Median : 0.0000 Median : 0.0000
## Mean : 0.0002447 Mean : 0.1111 Mean : 0.1286
## 3rd Qu.: 0.0000000 3rd Qu.: 0.0000 3rd Qu.: 0.0000
## Max. : 1.0000000 Max. : 1.0000 Max. : 1.0000
```

```
colSums(is.na(housing))
```


```
## longitude latitude
## 0 0
## housing_median_age total_rooms
## 0 0
## total_bedrooms population
## 0 0
## households median_income
## 0 0
## median_house_value `ocean_proximity<1H OCEAN`
## 0 0
## ocean_proximityINLAND ocean_proximityISLAND
## 0 0
## `ocean_proximityNEAR BAY` `ocean_proximityNEAR OCEAN`
## 0 0
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.3.3
```

```
## corrplot 0.95 loaded
```

```
corr_matrix <- cor(housing)
corrplot(corr_matrix, method = "color", tl.cex = 0.7)
```



```
set.seed(123) # ensures reproducibility

# Use caret's createDataPartition function
library(caret)

split <- createDataPartition(housing$median_house_value, p = 0.8, list = FALSE)

train_data <- housing[split, ]
test_data <- housing[-split, ]
```

0.0.0.5 Split the Data

```
nrow(train_data)
```

0.0.0.5.1 Check the Split

```
## [1] 16348
```

```
nrow(test_data)
```

```
## [1] 4085
```

```
colnames(train_data)
```

```
## [1] "longitude"          "latitude"
## [3] "housing_median_age" "total_rooms"
```

```
## [5] "total_bedrooms"      "population"
## [7] "households"          "median_income"
## [9] "median_house_value"  "`ocean_proximity<1H OCEAN`"
## [11] "ocean_proximityINLAND" "ocean_proximityISLAND"
## [13] "`ocean_proximityNEAR BAY`" "`ocean_proximityNEAR OCEAN`"
```

```
colnames(test_data)
```

```
## [1] "longitude"          "latitude"
## [3] "housing_median_age" "total_rooms"
## [5] "total_bedrooms"     "population"
## [7] "households"         "median_income"
## [9] "median_house_value" "`ocean_proximity<1H OCEAN`"
## [11] "ocean_proximityINLAND" "ocean_proximityISLAND"
## [13] "`ocean_proximityNEAR BAY`" "`ocean_proximityNEAR OCEAN`"
```

```
model_train <- lm(median_house_value ~ ., data = train_data)
summary(model_train)
```

0.0.0.6 Fit the Model on Training Data

```
##
## Call:
## lm(formula = median_house_value ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -560505  -42590  -10327   28586   734628
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -2.255e+06  9.861e+04  -22.870  < 2e-16 ***
## longitude       -2.672e+04  1.138e+03  -23.475  < 2e-16 ***
## latitude        -2.549e+04  1.122e+03  -22.713  < 2e-16 ***
## housing_median_age  1.074e+03  4.904e+01   21.911  < 2e-16 ***
## total_rooms      -6.607e+00  8.820e-01   -7.491  7.16e-14 ***
## total_bedrooms    1.088e+02  7.824e+00   13.913  < 2e-16 ***
## population       -3.597e+01  1.172e+00  -30.702  < 2e-16 ***
## households        3.785e+01  8.426e+00    4.492  7.09e-06 ***
## median_income     3.961e+04  3.828e+02  103.474  < 2e-16 ***
## ````ocean_proximity<1H OCEAN``  -4.668e+03  1.748e+03   -2.671  0.00758 **
## ocean_proximityINLAND    -4.443e+04  2.510e+03  -17.700  < 2e-16 ***
## ocean_proximityISLAND     1.477e+05  3.086e+04    4.786  1.71e-06 ***
## ````ocean_proximityNEAR BAY``  -7.389e+03  2.444e+03   -3.024  0.00250 **
## ````ocean_proximityNEAR OCEAN``      NA          NA        NA        NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 68870 on 16335 degrees of freedom
## Multiple R-squared:  0.6441, Adjusted R-squared:  0.6438
## F-statistic: 2463 on 12 and 16335 DF, p-value: < 2.2e-16
```

0.0.0.7 California Housing Price Prediction (Linear Regression)

```
predictions <- predict(model_train, newdata = test_data)
```

0.0.0.7.1 Predict on Test Data

```
# Calculate RMSE and R-squared
rmse <- sqrt(mean((test_data$median_house_value - predictions)^2))
r2 <- cor(test_data$median_house_value, predictions)^2

cat("RMSE:", rmse, "\nR-squared:", r2)
```

0.0.0.7.2 Evaluate Model Performance

```
## RMSE: 67873.12
## R-squared: 0.6558641
```


0.0.1 Regularization Techniques

```
x <- model.matrix(median_house_value ~ ., data = housing)[, -1] # remove intercept
y <- housing$median_house_value

set.seed(123)
train_idx <- sample(1:nrow(x), 0.8 * nrow(x))
x_train <- x[train_idx, ]
y_train <- y[train_idx]
x_test <- x[-train_idx, ]
y_test <- y[-train_idx]
```

0.0.1.1 Prepare Data for Regularization

```
ridge_model <- cv.glmnet(x_train, y_train, alpha = 0)
plot(ridge_model)
```



0.0.1.2 Ridge Regression (L2)

```
# Find best lambda (regularization parameter)
best_lambda_ridge <- ridge_model$lambda.min
best_lambda_ridge
```

0.0.1.2.1 regularization parameter

```
## [1] 7987.83
```

```
ridge_pred <- predict(ridge_model, s = ridge_model$lambda.min, newx = x_test)
```

0.0.1.3 Prediction_Ridge

```
ridge_rmse <- sqrt(mean((ridge_pred - y_test)^2))
ridge_r2 <- 1 - sum((ridge_pred - y_test)^2) / sum((y_test - mean(y_test))^2)
cat("Ridge Regression:\n")
```

0.0.1.4 Evaluation Ridgr Regression

```
## Ridge Regression:
```

```
cat(" RMSE:", ridge_rmse, "\n")
```


```
## RMSE: 68931.86
```



```
cat(" R-squared:", ridge_r2, "\n\n")
```

```
## R-squared: 0.63499
```

```
library(glmnet)
plot(ridge_model$glmnet.fit, xvar = "lambda")
```



0.0.1.4.1 Ridge Regression Diagnostics

```
coef(ridge_model, s = ridge_model$lambda.min)
```

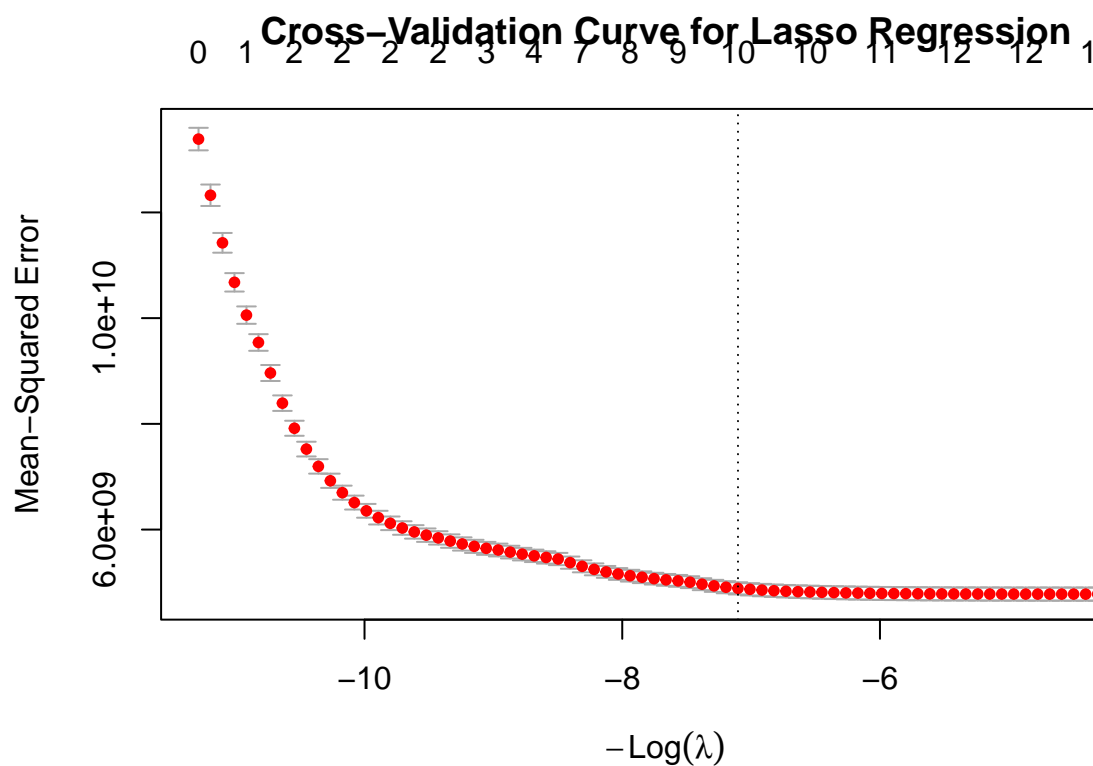
```
## 14 x 1 sparse Matrix of class "dgCMatrix"
## s=7987.83
## (Intercept) -7.958822e+05
## longitude -9.432991e+03
## latitude -8.406588e+03
## housing_median_age 1.003642e+03
## total_rooms 7.246974e-01
## total_bedrooms 4.134851e+01
## population -2.514532e+01
## households 4.091730e+01
## median_income 3.544553e+04
## ````ocean_proximity<1H OCEAN```` 1.770221e+04
## ocean_proximityINLAND -4.397253e+04
## ocean_proximityISLAND 1.912247e+05
## ````ocean_proximityNEAR BAY```` 2.023143e+04
## ````ocean_proximityNEAR OCEAN```` 2.758138e+04
```

```

# Perform Lasso Regression (L1 Regularization)
set.seed(123)
lasso_model <- cv.glmnet(x_train, y_train, alpha = 1)

# Cross-validated Lasso plot with title
plot(lasso_model, main = "Cross-Validation Curve for Lasso Regression")

```



0.0.1.5 Lasso Regression

```

# Best lambda
best_lambda_lasso <- lasso_model$lambda.min
best_lambda_lasso

```

```
## [1] 56.35256
```

```

# Lasso Regression
lasso_pred <- predict(lasso_model, s = lasso_model$lambda.min, newx = x_test)

```

0.0.1.6 Lasso Prediction

```

lasso_rmse <- sqrt(mean((lasso_pred - y_test)^2))
lasso_r2 <- 1 - sum((lasso_pred - y_test)^2) / sum((y_test - mean(y_test))^2)

cat("Lasso Regression:\n")

```

0.0.1.6.1 Evaluation

```
## Lasso Regression:
```

```
cat("  RMSE:", lasso_rmse, "\n")
```

```
##  RMSE: 67681.33
```

```
cat("  R-squared:", lasso_r2, "\n")
```

```
##  R-squared: 0.6481135
```

```
coef(lasso_model, s = best_lambda_lasso)
```

0.0.1.6.2 Lasso Regression Diagnostics

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##                                     s=56.35256
## (Intercept)                       -2.196103e+06
## longitude                         -2.596847e+04
## latitude                         -2.470979e+04
## housing_median_age                1.055381e+03
## total_rooms                      -5.562909e+00
## total_bedrooms                   8.775694e+01
## population                       -3.700647e+01
## households                        5.858007e+01
## median_income                    3.906495e+04
## `\\`ocean_proximity<1H OCEAN\\` ` .
## ocean_proximityINLAND            -3.921343e+04
## ocean_proximityISLAND             1.682914e+05
## `\\`ocean_proximityNEAR BAY\\` ` -2.839300e+03
## `\\`ocean_proximityNEAR OCEAN\\` ` 5.092122e+03
```

```
# Prepare data
```

```
x <- model.matrix(median_house_value ~ ., train_data)[, -1]
y <- train_data$median_house_value
```

```
# Create training control
```

```
train_control <- trainControl(method = "cv", number = 10)
```

```
# Fit PLS Regression model
```

```
set.seed(123)
```

```
pls_model <- train(
```

```
  x = x,
  y = y,
  method = "pls",
  trControl = train_control,
  tuneLength = 10,
  preProcess = c("center", "scale")
)
```

```
# Show best model
```

```
print(pls_model)
```

```
## Partial Least Squares
```

```
##
```

```
## 16348 samples
```

```
## 13 predictor
```

```
##
```

```

## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 14713, 14712, 14714, 14713, 14713, 14715, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      78045.73  0.5425952  58709.15
##   2      73743.99  0.5915625  54199.26
##   3      72147.29  0.6091188  52064.38
##   4      71400.04  0.6171795  51586.66
##   5      70456.33  0.6272779  51248.63
##   6      69672.94  0.6357131  50232.80
##   7      69545.56  0.6370352  50083.46
##   8      69322.88  0.6393446  49949.82
##   9      69100.78  0.6416603  49772.12
##  10      69072.99  0.6419112  49766.69
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 10.

# Evaluate on test data
x_test <- model.matrix(median_house_value ~ ., test_data)[, -1]
y_test <- test_data$median_house_value

pls_predictions <- predict(pls_model, newdata = x_test)


# Compute metrics
pls_rmse <- sqrt(mean((y_test - pls_predictions)^2))
pls_r2 <- cor(y_test, pls_predictions)^2

cat("PLS Regression:\n",
    "  RMSE:", pls_rmse, "\n",
    "  R-squared:", pls_r2, "\n")

## PLS Regression:
##   RMSE: 67902.41
##   R-squared: 0.6555516

library(caret)
plot(pls_model)

```



```
# If you used caret::train with method = "pls"
pls_coef <- coef(pls_model$finalModel, ncomp = 10) # ncomp = optimal number of components
pls_coef
```

```
## , , 10 comps
##
##
## .outcome
## longitude -41823.080
## latitude -42022.359
## housing_median_age 13738.501
## total_rooms -20290.447
## total_bedrooms 37074.890
## population -42468.432
## households 30392.935
## median_income 75905.282
## `\\`\\\\\\\\`ocean_proximity<1H OCEAN\\\\\\\\\\\\\\\\` 6732.612
## ocean_proximityINLAND -14691.680
## ocean_proximityISLAND 3375.103
## `\\`\\\\\\\\\\\\`ocean_proximityNEAR BAY\\\\\\\\\\\\\\\\` 3530.472
## `\\`\\\\\\\\\\\\\\\\`ocean_proximityNEAR OCEAN\\\\\\\\\\\\\\\\\\\\` 6948.294
```

```
set.seed(123)
ridge_cv <- cv.glmnet(x, y, alpha = 0, nfolds = 10)
best_lambda_ridge <- ridge_cv$lambda.min
rmse_ridge <- sqrt(min(ridge_cv$cvm)) # CV RMSE
```

```
# Lasso CV
```

```

lasso_cv <- cv.glmnet(x, y, alpha = 1, nfolds = 10)
best_lambda_lasso <- lasso_cv$lambda.min
rmse_lasso <- sqrt(min(lasso_cv$cvm)) # CV RMSE

# PLS CV
library(caret)
set.seed(123)
pls_model <- train(
  median_house_value ~ ., data = train_data,
  method = "pls",
  preProcess = c("center", "scale"),
  tuneLength = 10,
  trControl = trainControl(method = "cv", number = 10)
)
rmse_pls <- pls_model$results$RMSE[pls_model$results$ncomp == pls_model$bestTune$ncomp]

# Linear regression RMSE (CV)
set.seed(123)
lm_model <- train(
  median_house_value ~ ., data = train_data,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
)
rmse_lm <- lm_model$results$RMSE

# Combine results
cv_results <- data.frame(
  Model = c("Linear Regression", "Ridge Regression", "Lasso Regression", "PLS Regression"),
  CV_RMSE = c(rmse_lm, rmse_ridge, rmse_lasso, rmse_pls)
)

cv_results

##           Model  CV_RMSE
## 1 Linear Regression 68971.66
## 2 Ridge Regression 70126.12
## 3 Lasso Regression 69013.66
## 4 PLS Regression 69072.99

# Predictions on test set
x_test <- model.matrix(median_house_value ~ ., test_data)[, -1]
y_test <- test_data$median_house_value

# Linear Regression
pred_lm <- predict(lm_model, newdata = test_data)

# Ridge
pred_ridge <- predict(ridge_model, s = best_lambda_ridge, newx = x_test)

# Lasso
pred_lasso <- predict(lasso_model, s = best_lambda_lasso, newx = x_test)

# PLS
pred_pls <- predict(pls_model, newdata = test_data)

```

```

library(ggplot2)
plot_df <- data.frame(
  Actual = y_test,
  Linear = as.vector(pred_lm),
  Ridge = as.vector(pred_ridge),
  Lasso = as.vector(pred_lasso),
  PLS = as.vector(pred_pls)
)


library(tidyr)

## Warning: package 'tidyr' was built under R version 4.3.3
##
## Attaching package: 'tidyr'
## The following objects are masked from 'package:Matrix':
##
##     expand, pack, unpack
## The following object is masked from 'package:magrittr':
##
##     extract
# Convert to long format for ggplot
plot_long <- pivot_longer(plot_df, cols = -Actual, names_to = "Model", values_to = "Predicted")

ggplot(plot_long, aes(x = Actual, y = Predicted, color = Model)) +
  geom_point(alpha = 0.5) +
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "black") +
  labs(title = "Actual vs Predicted House Values",
       x = "Actual Median House Value",
       y = "Predicted Median House Value") +
  theme_minimal() +
  theme(legend.position = "bottom")


```

Actual vs Predicted House Values



```
ggplot(plot_long, aes(x = Actual, y = Predicted)) +  
  geom_point(alpha = 0.5, color = "steelblue") +  
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "red") +  
  labs(title = "Actual vs Predicted House Values", x = "Actual", y = "Predicted") +  
  facet_wrap(~Model) +  
  theme_minimal()
```



Actual vs Predicted House Values



```
# Predict on test data
y_pred <- predict(lasso_model, newx = x_test, s = "lambda.min")

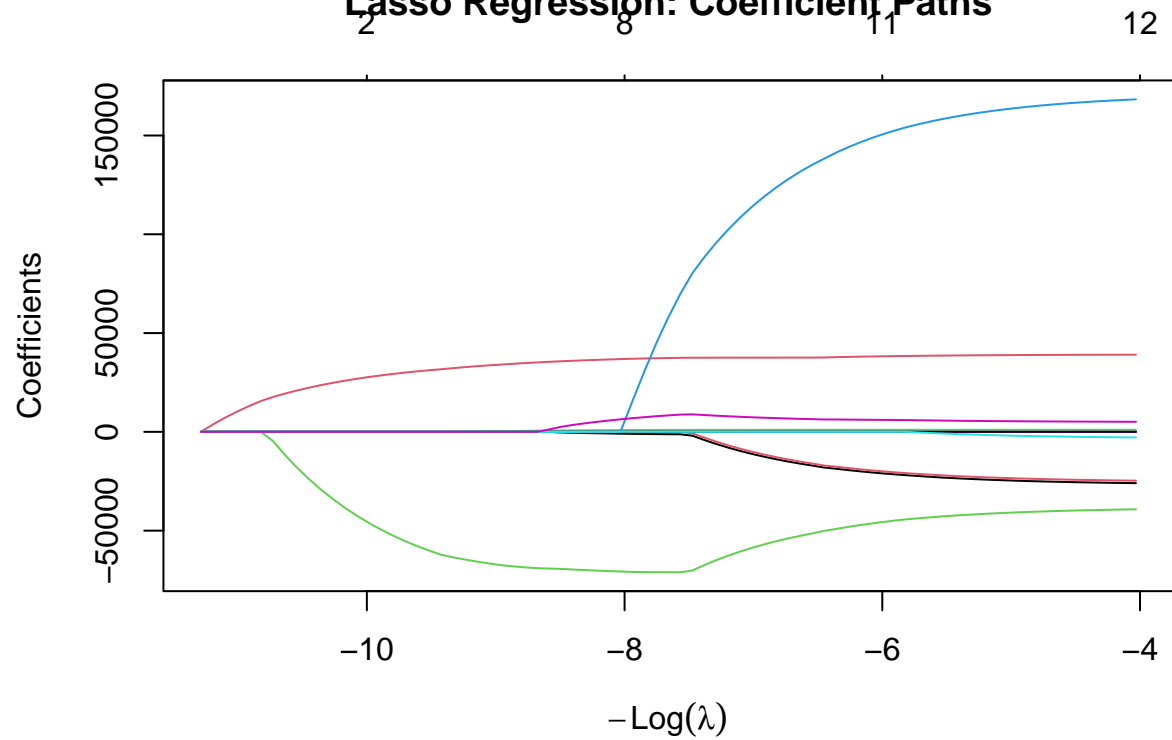
# Scatter plot: actual vs predicted
plot(y_test, y_pred,
     xlab = "Actual Median House Value",
     ylab = "Predicted Median House Value",
     main = "Lasso Regression: Actual vs Predicted",
     col = "blue", pch = 16)
abline(a = 0, b = 1, col = "red", lwd = 2) # perfect fit line
```

Lasso Regression: Actual vs Predicted



```
plot(lasso_model$glmnet.fit, xvar = "lambda", main = "Lasso Regression: Coefficient Paths")
```

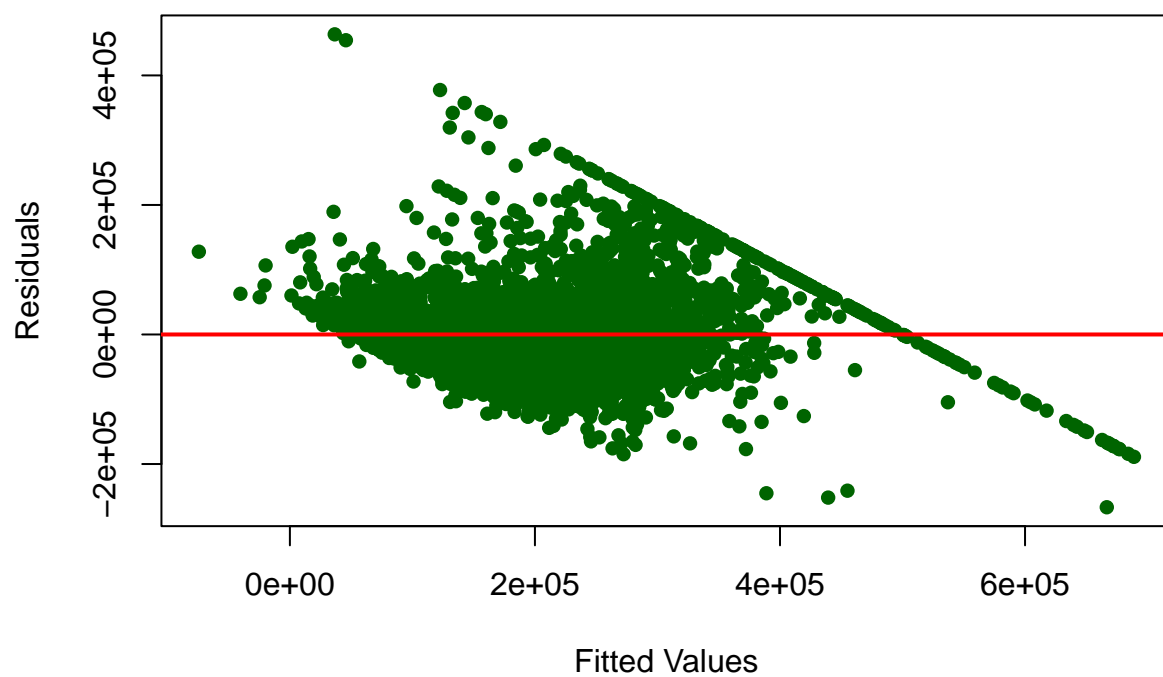
Lasso Regression: Coefficient Paths



```
# Residuals
residuals_lasso <- y_test - y_pred

# Residuals vs Fitted
plot(y_pred, residuals_lasso,
     xlab = "Fitted Values", ylab = "Residuals",
     main = "Lasso Residuals vs Fitted", pch = 16, col = "darkgreen")
abline(h = 0, col = "red", lwd = 2)
```

Lasso Residuals vs Fitted



```
# Normal Q-Q plot  
qqnorm(residuals_lasso, main = "Lasso Residual Q-Q Plot")  
qqline(residuals_lasso, col = "red", lwd = 2)
```

Lasso Residual Q-Q Plot

